

The banner is divided into three vertical panels. The left panel shows two men in a factory setting looking at a robotic arm. The middle panel shows a woman in a white lab coat sitting at a desk with a laptop. The right panel shows a woman in a server room holding a laptop. The Intel logo is in the top left, and the OpenVINO Toolkits text is in the top center.

intel[®] OpenVINO[™] Toolkits

Installation

visit openvino.ai

How to use

Examples




Powered by oneAPI

Three AI Engines in Intel® Core™ Ultra



NPU

Power Efficiency
Ideal for sustained AI workloads
and AI offload for battery life




CPU

Fast Response
Ideal for low-latency
AI workloads



GPU

High Throughput
Ideal for AI-accelerated digital
content creation and gaming



The right balance of power and performance for building and deploying AI models with OpenVINO™




Google Summer of Code



open source
initiative®

Apache 2.0

OpenVINO™



Powered
by oneAPI




OpenVINO™

Optimized Performance

CPU



GPU



NPU



FPGA



Windows

Linux

macOS


Benefits of Maximizing AI Performance




Dynamic
Model Loading



Efficient Resource
Management



High-Performance
Inference



User-Friendly
API




Installation

Installation



Installation

```
pip install openvino
```



OpenVINO GenAI Installation



Install OpenVINO™ 2025.4 #


Package	OpenVINO Base Package	OpenVINO GenAI	OpenVINO Model Server			
Version	2025.4 Recommended	2024.6	Nightly Build			
Operating System	Windows	macOS	Linux			
Distribution	PIP Includes NPU plugin Python API only	Archives Includes NPU plugin	npm JavaScript API only	GitHub Source	Gitee Source	Docker

```
pip install openvino-genai
```

The background image shows a group of people in a professional setting, possibly a meeting or collaborative workspace. One person is holding a large sheet of paper with various charts and graphs, while another is using a tablet. The scene is overlaid with a semi-transparent blue rectangle. Decorative elements include several small squares in shades of blue and green scattered across the image.

How to use

Quickstart



```
pip install openvino
```

```
import openvino as ov

img = load_img()

core = ov.Core()

# Tensorflow, ONNX, Keras, Tensorflow Lite, Paddlepaddle
compiled_model = core.compile_model(model="model.onnx", device_name="CPU")

output_layer = compiled_model.outputs[0]

result = compiled_model(img)[output_layer]
```

Supported Devices




CPU

```
compiled_model = core.compile_model(model=model, device_name="CPU")
```




Supported Devices



```
compiled_model = core.compile_model(model=model, device_name="GPU")
```




Supported Devices




```
compiled_model = core.compile_model(model=model, device_name="NPU")
```



GenAI Pipeline



How to get an OpenVINO GenAI model



AI & ML interests


Deep Learning Inference, Deep Learning Model Optimization

Recent Activity

- tybulewicz updated a model about 12 hours ago
[OpenVINO/efficientnet_b1_ft_in1k-fp16-ov](#)
- amokrov published a model 1 day ago
[OpenVINO/gemma-4-31B-it-int8-ov](#)
- amokrov updated a model 1 day ago
[OpenVINO/gemma-4-31B-it-int8-ov](#)


[View all activity](#)

Team members 274



[OpenVINO's models](#) 355 [Sort: Recently updated](#)

- | | |
|--|--|
| OpenVINO/efficientnet_b1_ft_in1k-fp16-ov
Image Classification • Updated about 12 hours ago | OpenVINO/gemma-4-31B-it-int8-ov
Image-Text-to-Text • Updated 1 day ago • 17 |
| OpenVINO/gemma-4-31B-it-int4-ov
Image-Text-to-Text • Updated 1 day ago • 8 | OpenVINO/gemma-4-E4B-it-int8-ov
Image-Text-to-Text • Updated 1 day ago • 2 |
| OpenVINO/gemma-4-E4B-it-int4-ov
Image-Text-to-Text • Updated 1 day ago • 14 | OpenVINO/efficientnet_b7-fp16-ov
Image Classification • Updated 6 days ago |
| OpenVINO/resnet50-fp16-ov
Image Classification • Updated 6 days ago | OpenVINO/efficientnet_v2_m-int8-ov
Image Classification • Updated 6 days ago |
| OpenVINO/resnet152-fp16-ov
Image Classification • Updated 6 days ago | OpenVINO/resnet152-int8-ov
Image Classification • Updated 6 days ago |
| OpenVINO/resnet50-int8-ov
Image Classification • Updated 6 days ago | OpenVINO/resnet18-int8-ov
Image Classification • Updated 6 days ago |
| OpenVINO/resnet18-fp16-ov
Image Classification • Updated 6 days ago | OpenVINO/resnet101-int8-ov
Image Classification • Updated 6 days ago |
| OpenVINO/resnet101-fp16-ov
Image Classification • Updated 6 days ago | OpenVINO/mobilenet_v3_small-int8-ov
Image Classification • Updated 6 days ago |
| OpenVINO/mobilenet_v3_small-fp16-ov
Image Classification • Updated 6 days ago | OpenVINO/mobilenet_v3_large-int8-ov
Image Classification • Updated 6 days ago |
| OpenVINO/mobilenet_v3_large-fp16-ov
Image Classification • Updated 6 days ago | OpenVINO/mobilenet_v2-int8-ov
Image Classification • Updated 6 days ago |



OpenVINO官方 Model Center

🔗 链接

🗑️ 删除组织

您是该组织的管理员，页面数据相关内容仅您可见

🏠 <https://github.com/openvinotoolkit/openvino>

关于我们

OpenVINO官方Model Center 是OpenVINO官方大模型的集散地，汇集适配Intel硬件的国内外最新大模型集合。为方便社区开发者积极参与，Model Center提供社区模型合集。• Community Models: 供社区开发者自行转换模型 • Community Premium Collection: 已被开发者验证的社区贡献模型 模型转换指南，请参考文档 <https://www.modelscope.cn/learn/2401> 欢迎开发者们积极使用、转换模型！

组织数据

组织成员

您是该组织的管理员



全部 773 合集 8 模型 416 研习社 348 品牌馆 1

合集

在当前组织搜索

最近更新

Speech-to-Text

@ OpenVINO官方 Model Ce... 34 1

whisper-large-v3-int8-ov
@OpenVINO官方 Model Center 1.1k 0

openai-whisper-large-v3-fp16-ov
@OpenVINO官方 Model Center 856 0

whisper-base-int4-ov
@OpenVINO官方 Model Center 1.1k 0

Speculative Decoding Draft Models

@ OpenVINO官方 Model Ce... 2 1

Llama-3.1-8B-Instruct-FastDraft-150M-int8-ov
@OpenVINO官方 Model Center 1.1k 0

Phi-3-mini-FastDraft-50M-int8-ov
@OpenVINO官方 Model Center 1.1k 0

暂无更多内容

Community Models

@ OpenVINO官方 Model Ce... 72 2

供OpenVINO社区贡献模型收录合集

PaddleOCR-VL-1.5-ov
@zhaohb 2.6k 3

Qwen3-VL-Embedding-2B-fp16-ov
@rainhenry 135 0

HY-MT1.5-1.8B-nf4-ov-npu-r1p8
@rainhenry 6 0

LLM

@ OpenVINO官方 Model Ce... 140 3

bloomz-1b1-fp16-ov
@OpenVINO官方 Model Center 907 1

gemma-7b-it-int8-ov
@OpenVINO官方 Model Center 1.1k 0

gemma-2b-fp16-ov
@OpenVINO官方 Model Center 944 0

共8个合集 查看全部

模型

在当前组织搜索

最近更新

OpenVINO/Phi-3.5-mini-instruct-int4-gq-ov NEW

OpenVINO PyTorch mit ...

OpenVINO官方 Model Center 2026.04.22 29

OpenVINO/Qwen3-pruned-6L-from-0.6B-int8-ov

OpenVINO PyTorch other ...

OpenVINO官方 Model Center 2026.04.21 246

OpenVINO/gemma-4-31B-it-int4-ov NEW

视觉多模态理解 Transformers等3个框架 apache-2.0 gemma4 conversational

OpenVINO官方 Model Center 2026.04.21 3

OpenVINO/gemma-4-E4B-it-int4-ov NEW

视觉多模态理解 Transformers等3个框架 apache-2.0 gemma4 conversational

OpenVINO官方 Model Center 2026.04.21 3

OpenVINO/Phi-3-mini-FastDraft-50M-int8-sym-ov

OpenVINO PyTorch other ...

OpenVINO官方 Model Center 2026.04.21 56

OpenVINO/efficientnet_b1_ft_int1k_fp16-ov NEW

通用分类 PyTorch apache-2.0 Classification image-classification vision

OpenVINO官方 Model Center 2026.04.22 3

OpenVINO/gemma-4-31B-it-int8-ov NEW

视觉多模态理解 Transformers等3个框架 apache-2.0 gemma4 conversational

OpenVINO官方 Model Center 2026.04.21 3

OpenVINO/gemma-4-E4B-it-int8-ov NEW

视觉多模态理解 Transformers等3个框架 apache-2.0 gemma4 conversational

OpenVINO官方 Model Center 2026.04.21 5

OpenVINO/Phi-3-mini-4k-instruct-int4-gq-ov NEW

OpenVINO PyTorch mit ...


OpenVINO官方 Model Center 2026.04.21 20

OpenVINO/mobilenet_v2-fp16-ov NEW

通用分类 PyTorch bsd-3-clause image-classification vision

OpenVINO官方 Model Center 2026.04.17 3


共416个模型 查看全部



OpenVINO GenAI model workflow


Export and optimize pipeline

```
optimum-cli export openvino --model <model_id> --weight-format  
fp16 <exported_model_name>
```





Deploy pipeline

```
import openvino_genai as ov_genai  
  
pipe = ov_genai.LLMPipeline(model_path, "GPU")  
print(pipe.generate("What is OpenVINO?", max_length=200))
```



OpenVINO GenAI Pipeline



Text Generation




```
import openvino_genai as ov_genai

pipe = ov_genai.LLMPipeline(model_path, "GPU")
print(pipe.generate("What is OpenVINO?", max_length=200))
```




```
#include "openvino/genai/llm_pipeline.hpp"
#include <iostream>

int main(int argc, char* argv[]) {
    ov::genai::LLMPipeline pipe(model_path, "GPU");
    std::cout << pipe.generate("What is Large Language
Model?", ov::genai::max_new_tokens(200));
}
```



Text Generation

Chat History



```
import openvino_genai

config = openvino_genai.GenerationConfig()
config.max_new_tokens = 100

chat_history = openvino_genai.ChatHistory()

while True:
    prompt = input("question:\n")
    chat_history.append({"role": "user", "content": prompt})

    decoded_results = pipe.generate(chat_history, config, streamer)
    chat_history.append({"role": "assistant", "content": decoded_results.texts[0]})
```

Text Embedding



```
import openvino_genai as ov_genai


model_path=("BAAI/bge-small-en-v1.5")

config = openvino_genai.TextEmbeddingPipeline.Config()
pipeline = ov_genai.TextEmbeddingPipeline(model_path, "GPU", config)
embeddings = pipeline.embed_documents(["What is OpenVINO?"])
```




```
#include "openvino/genai/rag/text_embedding_pipeline.hpp"
#include <iostream>

int main(int argc, char* argv[]) {
    ov::genai::TextEmbeddingPipeline::Config config;
    ov::genai::TextEmbeddingPipeline pipe(model_path, "GPU", config);
    ov::genai::EmbeddingResults embeddings =
        pipeline.embed_documents(documents);
}
```



Text Reranking



```
import openvino_genai as ov_genai

config = ov_genai.TextRerankPipeline.Config()
config.top_n = 3


pipeline = ov_genai.TextRerankPipeline(model_dir, "GPU", config)
rerank_result = pipeline.rerank(query, texts)
```



```
#include "openvino/genai/text_rerank_pipeline.hpp"
#include <iostream>

int main(int argc, char* argv[]) {
    ov::genai::TextRerankPipeline::Config config;
    config.top_n = 3;

    ov::genai::TextRerankPipeline pipeline(models_path, "GPU", config);
    std::vector<std::pair<size_t, float>> rerank_result =
        pipeline.rerank(query, documents);
}
```



Text-to-Image generation

Demo Text to Image with OpenVINO with Generate API

Anime, masterpiece, high quality, a green snowman with a happy: Run



☐ Use LoRA

Advanced Settings ◀

☰ Examples

Cyberpunk cityscape like Tokyo New York with tall buildings at dusk golden hour cinematic lighting

Curly-haired unicorn in the forest, anime, line

Anime, masterpiece, high quality, a green snowman with a happy smiling face in the snows

A panda reading a book in a lush forest.

Pirate ship sailing on a sea with the milky way galaxy in the sky and purple glow lights




```
import openvino_genai as ov_genai
pipe = openvino_genai.Text2ImagePipeline(model_dir, "GPU")
image_tensor = pipe.generate(prompt)
```




```
#include "openvino/genai/image_generation/text2image_pipeline.hpp"

int main(int argc, char* argv[]) {
    ov::genai::Image2ImagePipeline pipe(models_path, "GPU");
    ov::Tensor image = pipe.generate(prompt, image);
}
```


Text-to-image generation with LoRA



Without LoRA



With LoRA (Alpha=1)



```
import openvino_genai as ov_genai


pipe = openvino_genai.Text2ImagePipeline(model_dir, "GPU")
image_tensor = pipe.generate(prompt)
```



```
#include "openvino/genai/image_generation/text2image_pipeline.hpp"

int main(int argc, char* argv[]) {
    ov::genai::Text2ImagePipeline pipe(models_path, "GPU");
    ov::Tensor image = pipe.generate(prompt);
}
```

Image-to-image Generation



```
import openvino_genai as ov_genai




pipe = openvino_genai.Image2ImagePipeline(model_dir, "GPU")
image_tensor = pipe.generate(prompt, image, strength=0.8)
```



```
#include "openvino/genai/image_generation/image2image_pipeline.hpp"

int main(int argc, char* argv[]) {
    ov::genai::Image2ImagePipeline pipe(models_path, "GPU");
    ov::Tensor image = pipe.generate(prompt, image,
                                     ov::genai::strength(0.8f));
}
```

Inpainting



```
import openvino_genai as ov_genai


pipe = openvino_genai.InpaintingPipeline(model_dir, "GPU")
image_tensor = pipe.generate(prompt, image, mask_image)
```

```
#include "openvino/genai/image_generation/text2image_pipeline.hpp"

int main(int argc, char* argv[]) {
    ov::genai::InpaintingPipeline pipe(models_path, "GPU");
    ov::Tensor generated_image = pipe.generate(prompt, image, mask_image)
}
```



Image processing with Large Visual Models



```
import openvino_genai as ov_genai

pipe = ov_genai.VLMPipeline(model_dir, "GPU")
# read images to OV Tensors
rgbImages = read_images("filepath")
print(pipe.generate(prompt, images=rgbImages, max_new_tokens=100))
```

```
int main(int argc, char* argv[]) {
    ov::genai::VLMPipeline pipe(model_dir, "GPU");
    // read images into vector of OV Tensors
    std::vector<ov::Tensor> rgbImages = read_images("filepath");
    std::cout << pipe.generate(prompt, ov::genai::images(rgbImages),
                                ov::genai::max_new_tokens(100));
}
```



Speech transcription and translation



```
Windows PowerShell
(venv_genai) PS C:\Users\user\genai_demos\openvino.genai\samples\python\whisper_speech_recognition> python .\recorder.py
```



```
import openvino_genai as ov_genai


pipe = openvino_genai.WhisperPipeline(model_dir, "CPU")
# pipeline expects normalized audio with Sample Rate of 16kHz
raw_speech = read_wav("filepath")
result = pipe.generate(raw_speech)
```



```
#include "openvino/genai/whisper_pipeline.hpp"

int main(int argc, char* argv[]) {
    ov::genai::WhisperPipeline pipe(model_path, "CPU");
    // pipeline expects normalized audio with Sample Rate of 16kHz
    ov::genai::RawSpeechInput raw_speech = read_wav("filepath");
    auto result = pipe.generate(raw_speech, config);
}
```

Speech generation



```
import openvino_genai as ov_genai
import soundfile as sf

pipe = openvino_genai.Text2SpeechPipeline(model_dir, "CPU")
result = pipe.generate("Hello! This is a speech generated by OpenVINO
                        Gen AI toolkit")
sf.write(output_file_name, result.speeches[0].data[0], samplerate=16000)
```



```
#include "openvino/genai/speech_generation/text2speech_pipeline.hpp"

int main(int argc, char* argv[]) {
    ov::genai::Text2SpeechPipeline pipe(model_path, "CPU");
    auto result = pipe.generate("What is OpenVINO?")[0];
}
```


Video generation



```
Loading pipeline components ... 100%
'loss_type=None' was set in the config but it is unrecognized. Using the default loss: 'ForCausalLMLoss'.
/home2/anisha/ltx_video_demo/.venv/lib/python3.10/site-packages/diffusers/models/attention.py:471: TracerWarning: Converting a tensor to a Python boolean might
rd the data flow of Python values, so this value will be treated as a constant in the future. This means that the trace might not generalize to other inputs!
  if current_length != target_length:
/home2/anisha/ltx_video_demo/.venv/lib/python3.10/site-packages/diffusers/models/attention.py:486: TracerWarning: Converting a tensor to a Python boolean might
rd the data flow of Python values, so this value will be treated as a constant in the future. This means that the trace might not generalize to other inputs!
  if attention_mask.shape[0] < batch_size * head_size:
INFO:nnf:Statistics of the bitwidth distribution:


| Weight compression mode | % all parameters (layers) | % ratio-defining parameters (layers) |
|-------------------------|---------------------------|--------------------------------------|
| int8_asym, per-channel  | 100% (287 / 287)          | 100% (287 / 287)                     |


Applying Weight Compression
INFO:nnf:Statistics of the bitwidth distribution:


| Weight compression mode | % all parameters (layers) | % ratio-defining parameters (layers) |
|-------------------------|---------------------------|--------------------------------------|
| int8_asym, per-channel  | 100% (45 / 45)            | 100% (45 / 45)                       |


Applying Weight Compression
INFO:nnf:Statistics of the bitwidth distribution:


| Weight compression mode | % all parameters (layers) | % ratio-defining parameters (layers) |
|-------------------------|---------------------------|--------------------------------------|
| int8_asym, per-channel  | 100% (45 / 45)            | 100% (45 / 45)                       |


Applying Weight Compression
INFO:nnf:Statistics of the bitwidth distribution:


| Weight compression mode | % all parameters (layers) | % ratio-defining parameters (layers) |
|-------------------------|---------------------------|--------------------------------------|
| int8_asym, per-channel  | 100% (170 / 170)          | 100% (170 / 170)                     |


Applying Weight Compression
>> OK: INT8 exported to /home2/anisha/ltx_video_demo/LTX-Video/INT8
=====
Installation complete
=====
FP16: /home2/anisha/ltx_video_demo/LTX-Video/FP16
INT8: /home2/anisha/ltx_video_demo/LTX-Video/INT8 (ready)

Run: bash run_inference.sh ["prompt"]
INT8: bash run_inference.sh --int8 ["prompt"]
=====
anisha@4thSPR-Flex:/home2/anisha/ltx_video_demo$
```



```
import openvino_genai as ov_genai
```

```
pipe = ov_genai.Text2VideoPipeline(model_dir, "CPU")
result = pipe.generate("A cat walking along a cobblestone").video
save_video("genai_video.avi", result)
```

```
#include "openvino/genai/video_generation/text2video_pipeline.hpp"
#include "imwrite_video.hpp"
```


```
int main(int argc, char* argv[]) {
    const std::string models_path = argv[1], prompt = argv[2];

    ov::genai::Text2VideoPipeline pipe(models_path, "CPU");
    ov::Tensor result = pipe.generate("A wooden boat").video;
    imwrite_video("genai_video.avi", result);
}
```



OpenVINO™ Model Server

Deployment Options: On-prem (Data Center)



Clients

Load Balancer


OpenVINO™ Model Server Containers

Available endpoints:

- Predict
- Chat completions
- Embeddings
- Reranking
- Image (VLM)
- Image generation

OpenVINO™ Model Server

Quick Start – Deploy



Docker with CPU only

```
docker run -d -u $(id -u) --rm -v ${PWD}/model:/model -p 9000:9000
openvino/model_server:latest --model_name faster_rcnn --model_path /model --port 9000
```

Docker with CPU and GPU


```
docker run --user $(id -u):$(id -g) -d --device /dev/dri --group-add=$(stat -c "%g"
/dev/dri/render*) --rm -p 8000:8000 -v $(pwd)/models:/models:rw
openvino/model_server:latest-gpu --source_model OpenVINO/Phi-3.5-mini-instruct-int4-ov -
-model_repository_path models --rest_port 8000 --target_device GPU --cache_size 2
```

Bare metal

```
ovms --source_model OpenVINO/Phi-3.5-mini-instruct-int4-ov --model_repository_path
models --rest_port 8000 --target_device GPU --cache_size 2
```

OpenVINO™ Model Server


Chat completions – CLI Client



```
curl http://localhost:8000/v3/chat/completions \
-H "Content-Type: application/json" \
-d '{
  "model": "meta-llama/Meta-Llama-3-8B-Instruct",
  "max_tokens":30,
  "stream":false,
  "messages": [
    {
      "role": "system",
      "content": "You are a helpful assistant."
    },
    {
      "role": "user",
      "content": "What is OpenVINO?"
    }
  ]
}' | jq .
```

OpenVINO™ Model Server

Chat completions – Python Client



```
import requests
payload = {"model": " meta-llama/Meta-Llama-3-8B-Instruct",
           "messages": [ {"role": "user", "content": "Say this is a test" }]}
headers = {"Content-Type": "application/json", "Authorization": "not used"}
response = requests.post("http://localhost:8000/v3/chat/completions", json=payload, headers=headers)
print(response.text)
```

```
from openai import OpenAI

client = OpenAI(base_url="http://localhost:8000/v3", api_key="unused")

stream = client.chat.completions.create(
    model="meta-llama/Meta-Llama-3-8B-Instruct",
    messages=[{"role": "user", "content": "Say this is a test"}],
    stream=True,
)

for chunk in stream:
    if chunk.choices[0].delta.content is not None:
        print(chunk.choices[0].delta.content, end="", flush=True)
```

OpenVINO™ Model Server


Vision Language Models

```
import base64
from openai import OpenAI

def encode_image(image_path):
    with open(image_path, "rb") as image_file:
        return base64.b64encode(image_file.read()).decode("utf-8")

image = encode_image("/path/to/image")

client = OpenAI(base_url="http://localhost:8000/v3", api_key="unused")
response = client.chat.completions.create(model="openbmb/MiniCPM-V-2_6", stream=False,
    messages=[
        {
            "role": "user",
            "content": [
                {
                    "type": "text",
                    "text": "What is in this image?",
                },
                {
                    "type": "image_url",
                    "image_url": {"url": f"data:image/jpeg;base64,{image}"},
                },
            ],
        },
    ],
)
print(response.choices[0].message)
```



```
Windows PowerShell
PS D:\>
```

```
Windows PowerShell
(opencvino_env) PS D:\mcp_ovms>
```


The background is a blurred photograph of people in a meeting. A semi-transparent blue rectangle is centered over the image. Several small, semi-transparent squares in shades of blue and green are scattered around the blue rectangle.

Documents&Examples

docs.openvino.ai

OpenVINO

InstallationProduct PageBlogForumSupportPerformance

2026

Search

GET STARTED

Install OpenVINO

Learn OpenVINO

System Requirements

HOW TO USE - GENERATIVE AI WOR...

Generative Model Preparation

Inference with OpenVINO GenAI

Inference with Optimum Intel

OpenVINO Tokenizers

HOW TO USE - CONVENTIONAL AI ...

Conventional Model Preparation

Model Optimization - NNCF

Running Inference

Deployment on a Local System

PyTorch Deployment via "torch.compile"

HOW TO USE - MODEL SERVING

QuickStart - classic models

QuickStart - LLM models

Prepare a Model Repository

Deploy Model Server

Write a Client Application

Model Server Features

Performance tuning

Demos

Troubleshooting

REFERENCE DOCUMENTATION

API Reference

OpenVINO IR format and Operation ...

Compatibility and Support

Legacy Features

OpenVINO Extensibility

OpenVINO Security

ABOUT OPENVINO

OpenVINO 2026.1

OpenVINO Model Hub

See performance benchmarks for top AI models!

Explore now

OpenVINO is an open-source toolkit for deploying performant AI solutions in the cloud, on-prem, and on the edge alike. Develop your applications with both generative and conventional AI models, coming from the most popular model frameworks. Convert, optimize, and run inference utilizing the full potential of Intel® hardware. There are three main tools in OpenVINO to meet all your deployment needs:

OpenVINO GenAI

Run and deploy generative AI models

OpenVINO Base Package

Run and deploy conventional AI models

OpenVINO Model Server

Deploy both generative and conventional AI inference on a server

For a quick ramp-up, check out the [OpenVINO Toolkit Cheat Sheet \[PDF\]](#) and the [OpenVINO GenAI Quick-start Guide \[PDF\]](#)

PyTorch

TensorFlow

TensorFlowLite

PaddlePaddle

ONNX

Keras

OpenVINO

Optimized Performance

CPU

GPU

NPU

FPGA

intel ATOM

intel CORE

intel CORE ULTRA

intel XEON

arm

intel IRIS Xe MAX

intel ARC GRAPHICS

intel ARC GRAPHICS

intel ARC GPU

intel CORE ULTRA

FPGA AI Suite

OpenVINO

intel ai

github.com/openvinotoolkit/op
envino_notebooks

openvino_notebooksPublic

Edit PinsWatch50Fork717Star2k

latestBranchesTagsGo to fileAdd fileCode

eaidova

fix musicgen for transformers 4.40 (#2026) 8ae5a4a · 17 hours ago 1,614 Commits

.binder	specify runtime for binder (#1782)	2 months ago
.ci	Add method to push data to openvino_notebooks ovvp (#2...	2 days ago
.docker	Bump jinja2 from 3.1.3 to 3.1.4 in /.docker (#2000)	5 days ago
.github	Fix changed-files list on windows actions (#2022)	2 days ago
licensing	220-yolov5-accuracy-check-and-quantization demo (#592)	2 years ago
notebooks	fix musicgen for transformers 4.40 (#2026)	17 hours ago
selector	Notebooks status based on the treon nightly runs (#2008)	3 days ago
utils	using chat_template of llama3 (#1974)	3 weeks ago
.gitignore	Remove digits from directories and notebooks names (#1827)	2 months ago
CONTRIBUTING.md	update contributing guide with info how to run black (#2024)	19 hours ago
Dockerfile	update notebooks 2024.1 (#1963)	3 weeks ago
Jenkinsfile	Update Jenkinsfile	last year
LICENSE	Add notebooks and updated README from develop branch ...	3 years ago
Makefile	add makefile for CI (#348)	3 years ago
README.md	update notebooks 2024.1 (#1963)	3 weeks ago
README_cn.md	update links (#1867)	2 months ago
SECURITY.md	Add security.md (#1218)	10 months ago
check_install.py	remove framework specific deps from main requirements (#...	last month
requirements.txt	remove framework specific deps from main requirements (#...	last month

About

Jupyter notebook tutorials for OpenVINO™

machine-learningcomputer-visiondeep-learninginferenceopenvino

ReadmeApache-2.0 licenseSecurity policyActivityCustom properties2k stars50 watching717 forksReport repository

Releases

No releases publishedCreate a new release

Contributors111

+ 97 contributors

Deployments70

github-pages17 hours ago+ 69 deployments

Filter Categories by name

AI Trends

First Steps

API Overview

Convert

Optimize

Model Demos


Model Training

Live Demos

Notebooks 149

Filter notebooks by name


Sort: Recently Added



AI Trends Convert and Optimize YOLOv9 with OpenVINO™

Model Demos • Object Detection

View on GitHub



AI Trends Visual Question Answering and Image Captioning using BLIP and OpenVINO

Model Demos • Image-to-Text • Visual Question Answering


View on GitHub



AI Trends Zero-shot Image Classification with OpenAI CLIP and OpenVINO™

Model Demos • Zero-Shot Image Classification

View on GitHub



OpenVINO Tokenizers: Incorporate Text Processing Into OpenVINO Pipelines

First Steps • API Overview • Convert • Optimize • Text Classification • Text Generation

Resources


- <https://github.com/openvinotoolkit/openvino>
- <https://github.com/openvinotoolkit/openvino.genai>
- https://github.com/openvinotoolkit/openvino_notebooks
- <https://github.com/huggingface/optimum-intel>
- https://github.com/openvinotoolkit/model_server

下载HuggingFace模型的方法

方法一：设置镜像网站

```
pip install -U huggingface_hub  
export HF_ENDPOINT=https://hf-mirror.com
```

CLI



方法二：本地下载

```
git clone https://www.modelscope.cn/qwen/Qwen-1\_8B-Chat.git
```

CLI

```
optimum-cli export openvino --model Qwen-1\_8B-Chat ...
```

CLI



Q&A